

Morphosyntactic and Referential Cues to the Identification of Generic Statements

Phil Crone

pcrone@stanford.edu
Department of Linguistics
Stanford University

Michael C. Frank

mcf Frank@stanford.edu
Department of Psychology
Stanford University

Abstract

Generic sentences (e.g., “birds lay eggs”) express generalizations about kinds, as opposed specific individuals or groups of individuals (e.g., “all birds lay eggs”). Generics are an important method of transmitting cultural knowledge, but because there is no unique marker of genericity, identifying whether a sentence is generic is a challenge. Here we investigated how language users use morphosyntactic and pragmatic cues to determine whether naturalistic sentences should receive generic interpretations. Experiment 1 demonstrates the effect of the morphosyntactic features of a sentence’s subject noun phrase (NP) on generic interpretation. Experiments 2 and 3 reveal that when a sentence’s subject NP does not have an obvious reference in context, the sentence is more likely to receive a generic interpretation. These data suggest the beginnings of an account by which cues to genericity could be combined to make graded, contextual judgments about whether an utterance has an intended generic meaning.

Keywords: Psycholinguistics; pragmatics; generics.

Introduction

Generic sentences, which express generalizations about kinds rather than individuals, are an important route for the transmission of cultural knowledge (?). For example, the sentence “birds lay eggs” expresses a general property of the kind *bird*, whereas “all birds fly” means that every member of the set lays eggs. A key difference between generic and non-generic statements is that generics allow for exceptions: “birds lay eggs” is true despite the fact that at least half of birds do not. In contrast, “all birds lay eggs” is technically false, because male birds don’t (Prasada, 2000).

Generics are not consistently marked by any particular lexical, morphological, or syntactic convention, so how do we know that a sentence is generic? Prior work suggest that we use at least three types of cues to guide the interpretation of sentences as generic or non-generic: morphosyntactic features, pragmatic cues, and world knowledge. In English, the subject noun phrase (NP) of a generic sentence is often a bare plural (“birds fly”), but indefinite singular (“a bird has wings”) and definite singular (“the bird is a warm-blooded animal”). In contrast, definite plural NPs (“the birds have feathers”) are generally thought to force non-generic interpretations. Tense and aspect also cue whether a sentence is interpreted generically, such that simple present tense (“birds fly”) is typically more generic than e.g., present progressive (“birds are flying overhead”) or past (“birds flew past my window”) (Carlson, 1977; Krifka et al., 1995; Lyons, 1977).

Pragmatics and world knowledge can also influence whether a sentence is interpreted as generic or non-generic. For example, if a unique bird is present in the context of an utterance of a sentence with the subject NP “the bird,”

a non-generic interpretation in which this NP refers to the bird in context may be more likely. Conversely, if no such bird exists in the context, a generic interpretation may be preferred. Finally, world knowledge about the properties shared by members of a kind will influence the interpretation of potentially generic sentences. The sentence “a bird does not fly” is interpreted as a non-generic sentence about some particular bird (e.g., a penguin), given world knowledge that, in general, birds fly.

Adults and children as young as 3 show a preference for interpreting bare plurals as generic, as compared to definite plurals, and are more likely to interpret sentences as generic when the subject NP has no available referent in context (Gelman & Raman, 2003). By age 3, children are less likely to assign a generic interpretation to a sentence when its subject NP has a possible referent in the preceding linguistic context, and they can also use knowledge about whether properties are generalizable to kinds as evidence about genericity (Cimpian & Markman, 2008). Young children can also use the definiteness of subject NPs, as well as tense and aspect, in this process (Cimpian, Meltzer, & Markman, 2011).

The present study builds on the previous work in two ways. First, we focus on the quantitative relationship between factors in adults’ recognition of generics. The majority of previous work on the identification of generics has focused on children’s abilities, perhaps stemming in part from an assumption that children face an interpretive problem that is not as relevant for adults. Recent work emphasizing the probabilistic nature of language comprehension suggests that adults face a similar problem, however (Levy, 2008; Frank & Goodman, 2012). On this view, language users resolve uncertainty in comprehension via probabilistic inference to the most likely interpretation. In the specific case of identifying generics, we can view adults as reasoning about the likelihood of an utterance being generic given morphosyntactic features of the sentence, features of the context, and the listeners’ world knowledge. Second, we collect naturalistic examples of generic and non-generic sentences generated by study participants, allowing for a more realistic representation of how genericity is used in natural language.

Our plan is as follows. We begin by introducing an experimental paradigm that allows us to measure the relative genericity of the predicate statements that participants produce; this paradigm allows us to examine the role of morphosyntactic information in genericity (Experiment 1). We next investigate the interaction between morphosyntactic information and pragmatic information about reference (Experiments 2A and B). Finally, we identify a set of sentences

whose genericity is ambiguous and demonstrate the independence of morphosyntactic and pragmatic information (Experiment 3). Taken together, these experiments indicate that morphosyntactic and pragmatic cues are separable parts of a fundamentally graded inference about whether a sentence refers to individuals or kinds.

Experiment 1: Morphosyntactic Cues

As discussed above, the number and definiteness of a sentence’s subject NP influence its interpretation as generic or non-generic (Carlson, 1977; Krifka et al., 1995; Lyons, 1977). Previous work investigating morphosyntactic cues to genericity have fixed the number of the subject NP as either singular (Cimpian et al., 2011) or plural (Gelman & Raman, 2003) and only manipulated definiteness. In Experiment 1, we considered the effects of number, definiteness, and their interaction on the interpretation of generics. We asked participants to perform a sentence completion task in which the subject NP was provided. Participants were then asked to indicate whether the sentences they produced were about specific individuals or kinds.

Method

Participants We recruited 100 participants to participate through Amazon’s Mechanical Turk website. Participants were restricted to individuals within the United States and were paid 50 cents to complete the study. The study took approximately 14 minutes to complete. We excluded 4 participants for indicating that their native language was not English.

Stimuli Forty-eight nouns were chosen to use as the bases for subject NPs. To ensure diversity among these subject NPs, twenty-four nouns were animate and twenty-four were inanimate. For each study participant, a random set of twelve nouns were assigned morphosyntactic features using a 2×2 factorial design crossing number (singular, plural) with definiteness (definite, indefinite). Half of nouns assigned to each factorial point were animate and half were inanimate. Base nouns were then programmatically edited to reflect the assigned number and definiteness values and create full NPs. For example, if the noun “panda” were assigned values *plural* and *definite*, the full NP would be “the pandas.”

In the first part of the experiment, participants saw a single NP followed by a single-line text box. They were instructed to “write a sentence starting with the phrase below.” In the second part of the experiment, participants were shown the sentences they had written in the first part of the experiment. They were shown one sentence at a time and asked whether the sentence was about a specific *noun* (for singular NPs), a specific group of *nouns* (for plural NPs), or about *nouns* in general. Participants indicated their response using a 5-point Likert scale with the following values: “Definitely about a specific *noun*/group of *nouns* (=1),” “Probably about a specific *noun*/group of *nouns*,” “Not sure,” “Probably about *nouns* in general,” “Definitely about *nouns* in general (=5).”




Figure 1: Mean genericity ratings in Experiment 1 by definiteness and number. Error bars show 95% confidence intervals.

Procedure We first presented participants with four example noun phrases. After these first four items, participants were shown an example sentence that they could have used for the example noun phrase. These sentences were constructed to favor non-generic interpretations for all NP types. After seeing the examples, participants were informed that they would not receive any feedback for the rest of the experiment.

Participants then began the first part of the experiment. All forty-eight subject NPs were presented in pseudorandom order, counterbalanced so that no two consecutive NPs matched in both number and definiteness. We required participants to provide a sentence completion for each item and required that sentence completions be a minimum of six characters long. After completing part one, participants were told that they were entering the second part of the experiment and were informed that they would be evaluating sentences’ genericity. They then evaluated the sentences they had produced in part one in the manner described above. Once again, sentences were presented in a pseudorandom order counterbalanced so that no two consecutive NPs matched in both number and definiteness. After judging all forty-eight sentences, they were then required to provide their native language. We also measured reaction times for each item, measured from the time the item was presented until the time a response was submitted. In the second part of the experiment, we excluded responses whose reaction times were greater than 2 standard deviations from the mean.

Results and Discussion

We found that both definiteness and number affected mean genericity ratings (Figure 1). The definites that participants produced tended to be rated as less generic, while the indefinites that participants produced—especially “bare plural” indefinites—tended to be rated as more generic (see Table 1 for examples).

Table 1: Example Productions from Experiment 1. Generic sentences received ratings of 5; non-generics received ratings of 1.

Definiteness	Number	Genericity	Examples
Indefinite	Singular	Generic	“A cow eats grass.”, “A bicycle is a convenient form of transportation.”
Indefinite	Singular	Non-generic	“A dog is sleeping on the porch.”, “A light bulb was dropped and exploded.”
Indefinite	Plural	Generic	“Gorillas are primates.”, “Towels are useful after showering.”
Indefinite	Plural	Non-generic	“Cats are circling the fishtank.”, “Kites were flying at the beach.”
Definite	Singular	Generic	“The camel uses his humps to conserve water.”, “The clock tells time.”
Definite	Singular	Non-generic	“The bear is moving closer to us.”, “The bed is unmade.”
Definite	Plural	Generic	“The kangaroos carry babies in pouches.”, “The trumpets are loud.”
Definite	Plural	Non-generic	“The rabbits are digging holes in the yard.”, “The couches were dusty and old.”

To analyze these findings, we fit a linear mixed-effects model to predict participants’ ratings of their sentences for genericity. We examined the interaction between animacy, definiteness, and number of subject NPs.¹ The model shows main effects of animacy, definiteness, and number. Sentences with indefinite NP subjects were rated significantly more generic ($\beta = 1.745, t = 14.169, p < 0.001$). Sentences with singular NP subjects ($\beta = -0.302, t = -3.450, p < 0.001$), and inanimate subjects ($\beta = -0.503, t = -5.732, p < 0.001$), were rated less generic. In addition, the model revealed an interaction effect between definiteness and number such that indefinite singulars were rated significantly less generic ($\beta = -1.023, t = -9.913, p < 0.001$). Finally, there was a significant three-way interaction such that inanimate, indefinite singulars were rated significantly more generic ($\beta = 0.536, t = 3.669, p < 0.001$).

The results are consistent with previous findings that indefinite singulars and bare plurals facilitate generic interpretations compared to definite singulars and definite plurals, respectively (Cimpian et al., 2011; Gelman & Raman, 2003). However, these results also show that plurality is independently associated with genericity, which has not been previously acknowledged. The interaction between definiteness and plurality reveals a superadditive effect by which indefinite singulars support generic interpretation less than bare plurals. This is unsurprising, as bare plurals are often taken to be the canonical subject type for English generics (Carlson, 1977; Krifka et al., 1995; Lyons, 1977). The effect of animacy on genericity ratings is consistent with previous findings that both children and adults produce more generic statements in describing animals than in describing artifacts (Brandone & Gelman, 2009).

The most unexpected finding is the fact that definite plurals were found to be more generic overall than definite singulars, despite the general view that definite singulars, but not definite plurals, allow for generic interpretations in English. This

result had us consider whether our methodology measured some property other than genericity. However, inspection of definite plurals that received high genericity ratings suggested that these ratings were at least *prima facie* appropriate.

Experiment 2A: Contextual Cues in Production

Experiment 1 demonstrated the role of morphosyntactic and, in the case of animacy, world knowledge cues to identifying generic statements. However, this experiment left contextual cues unaddressed. Experiment 2A investigated the role of context as a cue to identifying generics. Specifically, this experiment examined whether language users would be less likely to treat a subject NP as generic when it has a possible referent in the context.

Method

Participants Details of recruitment were identical to Experiment 1; we excluded two participants for indicating that their native language was not English for a final sample of 98 participants.

Stimuli Stimuli were almost identical to those used in Experiment 1. However, in part 1 of the experiment, each subject NP was accompanied by an image depicting either one or five instances of the animal or artifact denoted by the subject NP (Figure 2). Images with only one instance of the animal or artifact were taken directly from the Bank of Standardized Stimuli version 2.0 (Brodeur, Guérard, & Bouras, 2014). Images with multiple instances of the animal or artifact were created from the single instance images. The number of individuals in the image either matched or mismatched the number of the subject NP. Half of items for each subject NP type were paired with matching images and half were paired with mismatching images. Participants were instructed to describe the image in their sentence completion. The images only appeared for the sentence completion portion of the experiment; they did not appear in the rating portion.

Procedure The procedure was identical to that of Experiment 1. For the four example items, two items were randomly chosen to have matching images, the other items had mismatching images.

¹All mixed-effects models were fit using the lme4 package in R version 3.1.2. The model specification was as follows: `response ~ animacy * definiteness * number + (animacy + definiteness + number | WorkerId) + (definiteness + number | subject)`. For each effect, p values were calculated by treating the t statistic as if it were a z statistic (Barr, Levy, Scheepers, & Tily, 2013).



Figure 2: Examples of images used in Experiments 2A, 2B, and 3. Clockwise from upper left: single shirt, multiple shirts, multiple peacocks, single peacock.




Figure 3: Mean genericity ratings from Experiment 2A (ratings by producers), plotted by picture/plurality match/mismatch, definiteness, and number. Error bars show 95% confidence intervals.

Results and Discussion

Genericity ratings are shown in Figure 3. We replicated the morphosyntactic effects in Experiment 1. In addition, there was a small but reliable effect of referential mismatch such that sentences generated to refer to a picture that mismatched their base NP in number were later rated as more generic.

As in Experiment 1, we fit a linear mixed-effects model to predict participants' genericity ratings of their sentences. We examined the interaction of animacy, definiteness, and number of the subject NP, as well as image match/mismatch.² The model revealed significant main effects of animacy, definiteness, and image. Sentences with indefinite subject NPs were rated more generic ($\beta = 1.023, t = 7.518, p < 0.001$), while sentences with inanimate subjects were rated less generic ($\beta = -0.742, t = -4.005, p < 0.001$). Sentences produced with an image that did not match the subject NP in num-

ber were also rated significantly more generic ($\beta = 0.389, t = 3.414, p < 0.001$). The model also revealed interaction effects between definiteness and number such that sentences with indefinite singular subjects were significantly less generic ($\beta = -0.340, t = -2.170, p < 0.05$), and between definiteness and image such that sentences with indefinite subjects that appears with mismatching images were less generic ($\beta = -0.426, t = -2.731, p < 0.001$). Finally, the model showed a significant three-way interaction such that inanimate, indefinite subjects with mismatching images were rated significantly more generic ($\beta = 0.567, t = 2.571, p < 0.05$). In general, these effects are similar to those in Experiment 1.

The most important finding of Experiment 2A was the result that referential mismatch increased genericity. This result is similar to the finding of Gelman and Raman (2003), who showed that children and adults are more likely to interpret pronouns as generic when the only possible contextual referent mismatches in number. However, Gelman and Raman hypothesized that this effect would only be seen for plural pronouns in a context with only singular referents. In contrast, we observed this (modest) effect in all NP conditions, suggesting that generics generally are supported when NPs fail to have a reference in context, regardless of whether these NPs are singular or plural. Nevertheless, this experiment focused on whether generic sentences are *produced* more often for a given NP when this NP has a contextual referent. This is a slightly different question, however, than whether *listeners* are more likely to interpret sentences as generic when the subject NP fails to refer in context. We investigated this question in Experiment 2B.

Experiment 2B: Contextual Cues in Comprehension

Method

Participants Recruitment details were as above, except the study took 4 minutes and compensation was 30 cents. We excluded five participants for non-English native languages; our final sample consisted of 185 participants.

Stimuli Each participant in Experiment 2B was assigned some set of sentences produced by a participant in Experiment 2A. This set consisted of the twenty-four sentences a participant in Experiment 2A produced while viewing a picture that matched the subject NP in number. Each participant in Experiment 2A was matched with two individuals in Experiment 2B. (In addition to the two non-native English speakers, three producers were excluded for producing offensive or inappropriate sentences).

Participants were told that other Mechanical Turk workers had produced the sentences they were evaluating in order to describe the same pictures they were seeing. Stimuli were presented in a manner similar to part two of Experiment 2A, but images were displayed above each sentence. Half of the sentences for each NP subject type were presented with an image that matched the subject NP in number, while half were

²The model specification was as follows: $\text{response} \sim \text{animacy} * \text{definiteness} * \text{number} * \text{image} + (\text{animacy} + \text{definiteness} + \text{number} + \text{image} \mid \text{WorkerId}) + (\text{definiteness} + \text{number} + \text{image} \mid \text{subject})$.




Figure 4: Mean genericity ratings from Experiment 2B (independent ratings), plotted by picture/plurality match/mismatch, definiteness, and number. Error bars show 95% confidence intervals.

presented with an image that mismatched in number. For each item that was seen with a matching image by one participant, that item was seen with a mismatching image by a second participant.

Procedure The procedure was the same as that for part 2 of Experiments 1 and 2A. Participants were instructed to pay attention to the images they saw and consider what the other Mechanical Turk worker was thinking when he/she produced each sentence. For the four example items, two items were randomly chosen to have matching images, the other items had mismatching images.

Results and Discussion

Findings were largely comparable to Experiments 1 and 2A: Listeners were more likely to interpret a sentence as generic when its subject NP failed to refer in context (Figure 4). This effect was consistent across all NP types examined in the study.

We fit an identical linear model to Experiment 2A. The model revealed significant main effects of definiteness, number, and image. Sentences with indefinite subject NPs were rated more generic ($\beta = 1.401, t = 11.718, p < 0.001$), while sentences with singular subjects were rated less generic ($\beta = -0.421, t = -3.642, p < 0.001$). Sentences produced with an image that did not match the subject NP in number were also rated significantly more generic ($\beta = 0.382, t = 3.604, p < 0.001$). The model also revealed interaction effects between definiteness and image such that sentences with indefinite subjects that appears with mismatching images were less generic ($\beta = -0.426, t = -2.731, p < 0.001$) and between animacy and image such that sentences with inanimate subjects and mismatching images were less generic ($\beta = -0.380, t = -2.508, p < 0.05$). Finally, the model showed a significant three-way interaction such that inanimate, indefinite subjects with mismatching images were rated signifi-

cantly more generic ($\beta = 0.431, t = 2.028, p < 0.05$).

Experiment 3: Ambiguous Sentences

Although the effect of the contextual referent mismatch was present in both Experiments 2A and 2B, the effect size was relatively small in both cases ($\beta = 0.389$ for Experiment 2A; $\beta = 0.382$ for Experiment 2B). We hypothesized that this was possibly due to the existence of large numbers of sentences that were unambiguously generic or non-generic. In Experiment 3, we investigated whether this effect would be larger if participants considered more ambiguous sentences. We gathered consistent ratings for the sentences from Experiment 1, then identified the most ambiguous of these sentences and performed the same referential mismatch manipulation as we used in Experiment 2.

Method

Participants In part one of this experiment, we recruited 110 participants through Mechanical Turk for 67 cents and excluded one participant for having a non-English native language for a total sample of 109; in part two, we recruited 100 participants and paid 20 cents, excluding five for a total sample of 95.

Stimuli Stimuli for first part of this experiment were drawn from the sentences created in Experiment 1. We chose a set of 1100 sentences generated in Experiment 1 that received genericity ratings between 2 and 4, inclusive. We then created 11 subsets of 100 randomly-chosen sentences from the original set. This process was repeated nine additional times to yield 110 sets of 100 sentences each. Because the ambiguous sentences from Experiment 1 were not equally distributed among different NP types, we could not ensure an equal distribution of NP types in these sets. Each participant was assigned one set of sentences, which were displayed without images as in the genericity judgment portion of Experiment 1.

The stimuli for the second part of this experiment were drawn from those used in the first part. We calculated the mean genericity rating given to each sentence by participants in part one. We then selected the 20 sentences for each NP type whose mean ratings were closest to 3 and which did not contain any obvious typographical errors. This yielded a set of 80 sentences, which were then divided into groups of 16 sentences with 4 sentences of each NP type. In constructing these sets, animacy was not controlled for and in general, each set contained more animate subjects than inanimate subjects. No set contained multiple sentences with the same subject. These sentences were presented with matching and mismatching images in the same manner as that of Experiment 2B.

Procedure In part one, participants first saw four example sentences with feedback, with the definite singular and bare plural sentences favoring generic interpretations. After seeing the example sentences, the target sentences were presented for rating in randomized order as in part 2 of Experiment 1.




Figure 5: Histograms of genericity ratings from the full set of sentences produced for Experiment 3, split by definiteness and number.




Figure 6: Mean genericity ratings from Experiment 3, plotted by picture/plurality match/mismatch, definiteness, and number. Error bars show 95% confidence intervals.

In part two, the procedure was identical to that of Experiment 2B.

Results and Discussion

The distribution of mean genericity ratings across sentences is shown in Figure 5. The majority of definites were not rated as being particular generic; in contrast the distribution was substantially more bimodal for indefinites. Our procedure resulted in the selection of sentences whose ranking was closest to three; these sentences were relatively rare for all participants.

Turning to the second part of the experiment, we again found a strikingly consistent effect of referential mismatch, though again this effect was small (Figure 6). As with the previous experiments, results from the second part of Experiment 3 were analyzed with linear mixed-effects models. We compared two models. The first predicted genericity rating from number, definiteness, image, and their interactions.

The second predicted genericity rating from image alone.³ There was no significant difference between the fits of the two models ($\chi^2(6) = 2.738, p > 0.1$), so we opted for the minimal model, which revealed that sentences presented with mismatching images were rated significantly more generic than those presented with matching images ($\beta = 0.326, t = 4.192, p < 0.001$).

General Discussion

Acknowledgments

Thanks to Daniel Lassiter, Rose Schneider, and the members of the Language and Cognition Lab. We gratefully acknowledge the support of ONR Grant N00014-13-1-0287.

References

- Barr, D. J., Levy, R., Scheepers, C., & Tily, H. J. (2013). Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*, 68(3), 255–278.
- Brandone, A. C., & Gelman, S. A. (2009). Differences in preschoolers' and adults' use of generics about novel animals and artifacts: A window onto a conceptual divide. *Cognition*, 110(1), 1–22.
- Brodeur, M. B., Guérard, K., & Bouras, M. (2014). Bank of standardized stimuli (BOSS) phase II: 930 new normative photos. *PLoS ONE*, 9(9), e106953.
- Carlson, G. N. (1977). *Reference to Kinds in English*. Ph.D. dissertation. University of Massachusetts, Amherst.
- Cimpian, A., & Markman, E. (2008). Preschool children's use of cues to generic meaning. *Cognition*, 107(1), 19–53.
- Cimpian, A., Meltzer, T. J., & Markman, E. (2011). Preschoolers' use of morphosyntactic cues to identify generic sentences: Indefinite singular noun phrases, tense, and aspect. *Child Development*, 82(5), 1561–1578.
- Frank, M. C., & Goodman, N. D. (2012). Predicting pragmatic reasoning in language games. *Science*, 336(6084), 998.
- Gelman, S. A. (2003). *The essential child: Origins of essentialism in everyday thought*. Oxford University Press.
- Gelman, S. A., & Raman, L. (2003). Preschool children use linguistic form class and pragmatic cues to interpret generics. *Child Development*, 74(1), 308–325.
- Krifka, M., Pelletier, F. J., Carlson, G. N., ter Meulen, A., Link, G., & Chierchia, G. (1995). Genericity: An introduction. In *The Generic Book*. Chicago: The University of Chicago Press.
- Levy, R. (2008). Expectation-based syntactic comprehension. *Cognition*, 106(3), 1126–1177.
- Lyons, J. (1977). *Semantics: Volume 1*. New York: Cambridge University Press.
- Prasada, S. (2000). Acquiring generic knowledge. *Trends in Cognitive Science*, 4(2), 66–72.

³The two model specifications were as follows: `response ~ definiteness * number * image + (image | WorkerId)` and `response ~ image + (image | WorkerId)`.